



**Department of Physics,
Computer Science & Engineering**

CPSC 410 – Operating Systems I

Process Description & Control

Keith Perkins

Adapted from original slides by Dr. Roberto A. Flores
Also from “CS 537 Introduction to Operating Systems” Arpaci-Dusseau

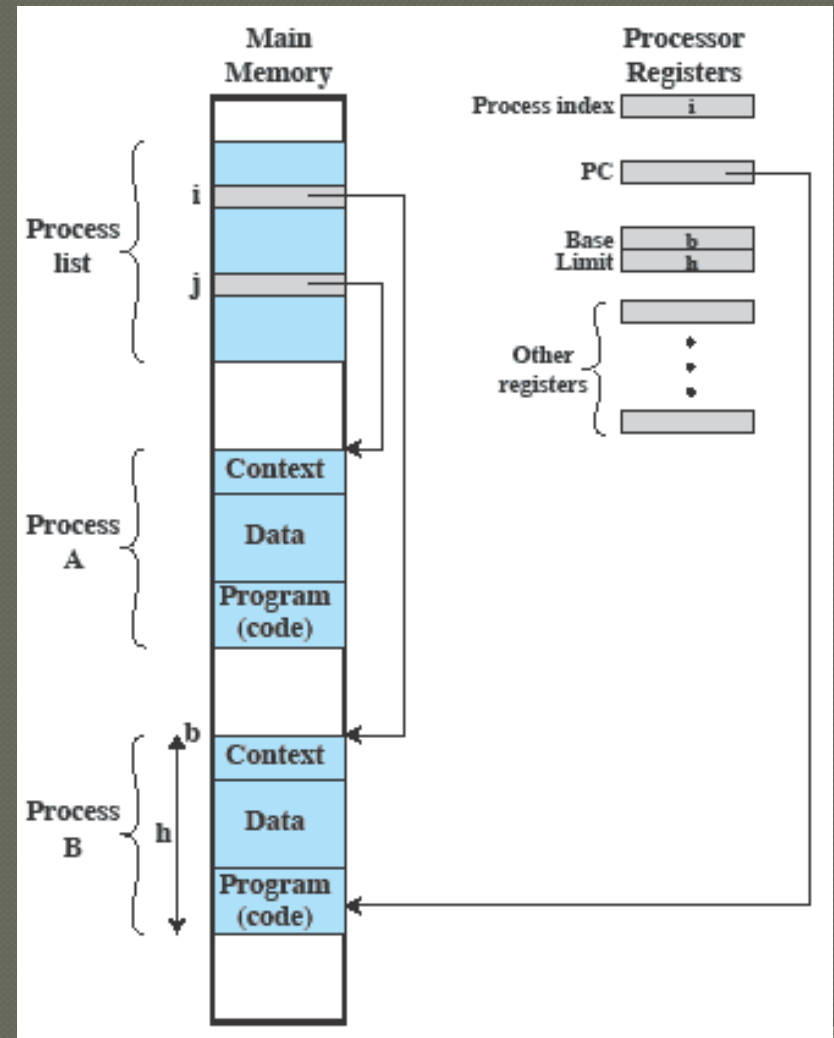
Topics

• Everything about Processes

- Control blocks
- States
- Description
- Control

Revisit - Process Management

- Scheduler chooses a process to run (more later)
- Dispatcher runs it
- How? What's in the Process List?
- BTW this list is a simplification



Processes

Control Blocks

- data structure created & managed by OS
 - **Identifier**: unique ID
 - **State**: (e.g., running, blocked)
 - **Priority**: relative to other processes
 - **Program counter**: address of next instruction
 - **Memory pointers**: to code & data
 - **I/O status**: I/O in use/pending
 - **Accounting**: CPU time used, IDs, ...
- data to hold/restore process state on interrupt/resume
 - key to support multiprocessing

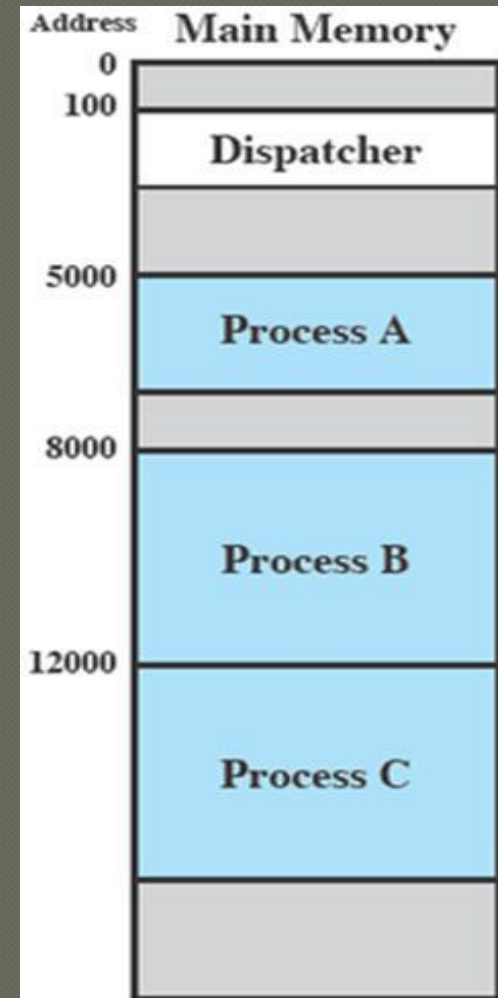
Identifier
State
Priority
Program counter
Memory pointers
Context data
I/O status information
Accounting information
⋮

Control blocks

States
Description
Control

- Dispatcher
 - Program that switches processes in/out of the CPU

Processes



Process dispatching mechanism

OS dispatching loop:

```
while(1) {
```

```
    run process for a while;
```

```
    save process state;
```

```
    next process = schedule (ready processes);
```

```
    load next process state;
```

```
}
```

Q1: how to gain control?



Q3: where to find processes?

Q2: what state must be saved?

Control blocks
States
Description
Control

Processes

States

- Trace
 - Instructions executed by a process
 - In multiprogramming:
 - interleaving of instructions as processes alternate using the CPU
- The pale blue lower right is dispatcher code
- Process switches because of Interrupts (timer, I/O)

5000	8000	12000
5001	8001	12001
5002	8002	12002
5003	8003	12003
5004		12004
5005		12005
5006		12006
5007		12007
5008		12008
5009		12009
5010		12010
5011		12011

(a) Trace of Process A (b) Trace of Process B (c) Trace of Process C

1	5000	27	12004
2	5001	28	12005
3	5002		
4	5003	29	100
5	5004	30	101
6	5005	31	102
		32	103
		33	104
		34	105
		35	5006
		36	5007
		37	5008
		38	5009
		39	5010
		40	5011
		41	100
		42	101
		43	102
		44	103
		45	104
		46	105
		47	12006
		48	12007
		49	12008
		50	12009
		51	12010
		52	12011

Timeout

Timeout

I/O Request

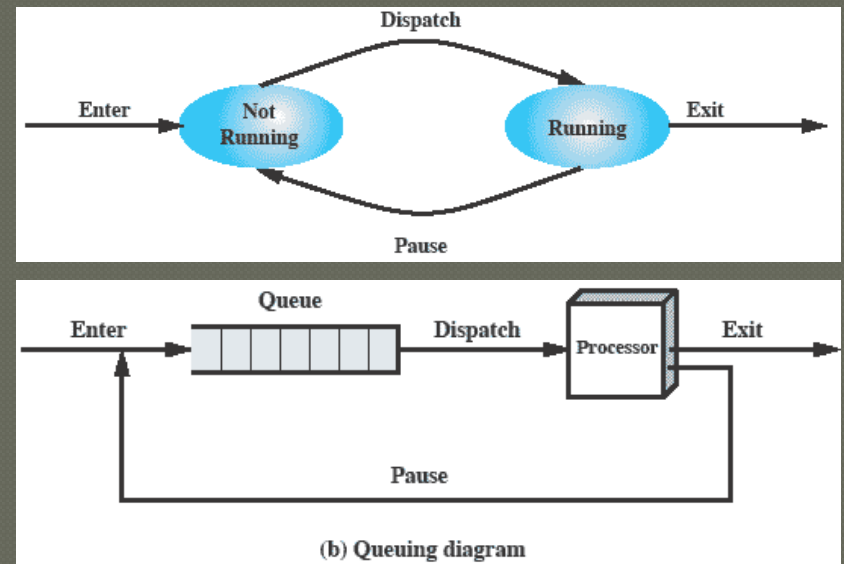
Timeout

Timeout

Processes

States (2 states)

- One CPU
- Round-robin (timeout)
- **Running**: CPU time!
- **Not running**: or not




- Where do processes come from?
- When do they stop?

Processes

- ◉ Where do processes come from? (start)
 - **Interactive logon**: User in terminal logs in
 - **OS service**: OS-provided service (e.g., print spooler)
 - **Spawned by process**: uses parallelism (parent spawns child)
- ◉ When do they end? (termination)
 - Normal
 - Job finishes, user logs off, OS shutting down, etc.
 - Abnormal
 - **Timeout**: timeslice
 - **Resource error**: out of memory, I/O device unresponsive, deadlock
 - **Runtime error**: arithmetic operation, uninitialized variable
 - **Authorization error**: memory out of bounds, resource/instruction privilege

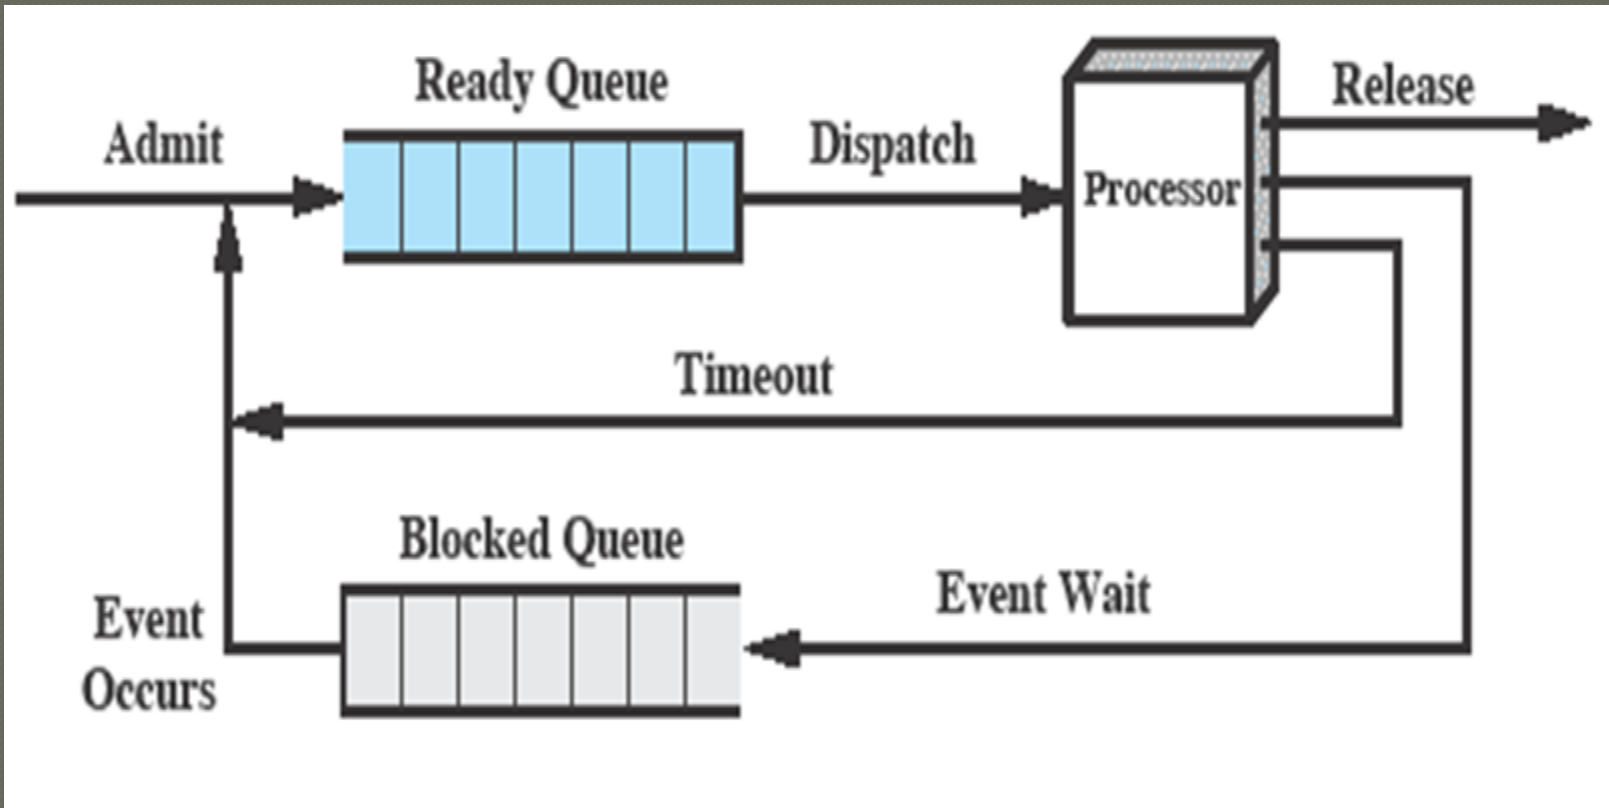
Processes

States (5 states)



- **New**: not yet in memory
- **Ready**: awaiting its turn
- **Running**: CPU time!
- **Blocked**: waiting for I/O
- **Exit**: done & gone

Using Two Queues





Control blocks
States
Description
Control

States (5 states)

- e.g., Processes A, B & C

- Multiple block queues (1 per I/O device)



Processes

States (7 states)

- What if running I/O intensive processes and all are waiting for I/O?
- **Solution: suspend blocked processes to disk, bring in new (from new or ready/suspend)**




● Process tables

- OS keeps list of processes. Each entry tracks data about each process (**process image**)
 - **Heap**:
 - **Globals**:
 - **Code**: program to execute
 - **stack**: method call stack frames
 - **process control block (PCB)**: data OS uses to control process
 - process **identification**: process/parent/user ID
 - processor **state information**: user/control registers, stack pointers
 - process **control information**: scheduling, inter-process comms, ...
- reference (directly/indirectly) memory, I/O & file tables

Processes

Process tables




● Process tables

Process identification

- Each process has a unique ID
 - IDs are used for reference:
 - in other tables
 - in inter-process communication
 - when a parent spawns a child process
- ➡ process **identification**: process/parent/user ID
 - ➡ processor **state information**: user/control registers, stack pointers
 - ➡ process **control information**: scheduling, inter-process comms, ...
 - reference to memory, I/O & file tables

Process state information

- stack pointers
- user-visible registers
- control & status registers
 - **program status word (PSW)**, e.g., EFLAGS in x86 processors



ID = Identification flag
 VIP = Virtual interrupt pending
 VIF = Virtual interrupt flag
 AC = Alignment check
 VM = Virtual 8086 mode
 RF = Resume flag
 NT = Nested task flag
 IOPL = I/O privilege level
 OF = Overflow flag


DF = Direction flag
 IF = Interrupt enable flag
 TF = Trap flag
 SF = Sign flag
 ZF = Zero flag
 AF = Auxiliary carry flag
 PF = Parity flag
 CF = Carry flag

- ➡ processor **state information**: user/control registers, stack pointers
- ➡ process **control information**: scheduling, inter-process comms, ...
- reference (directly/indirectly) memory, I/O & file tables

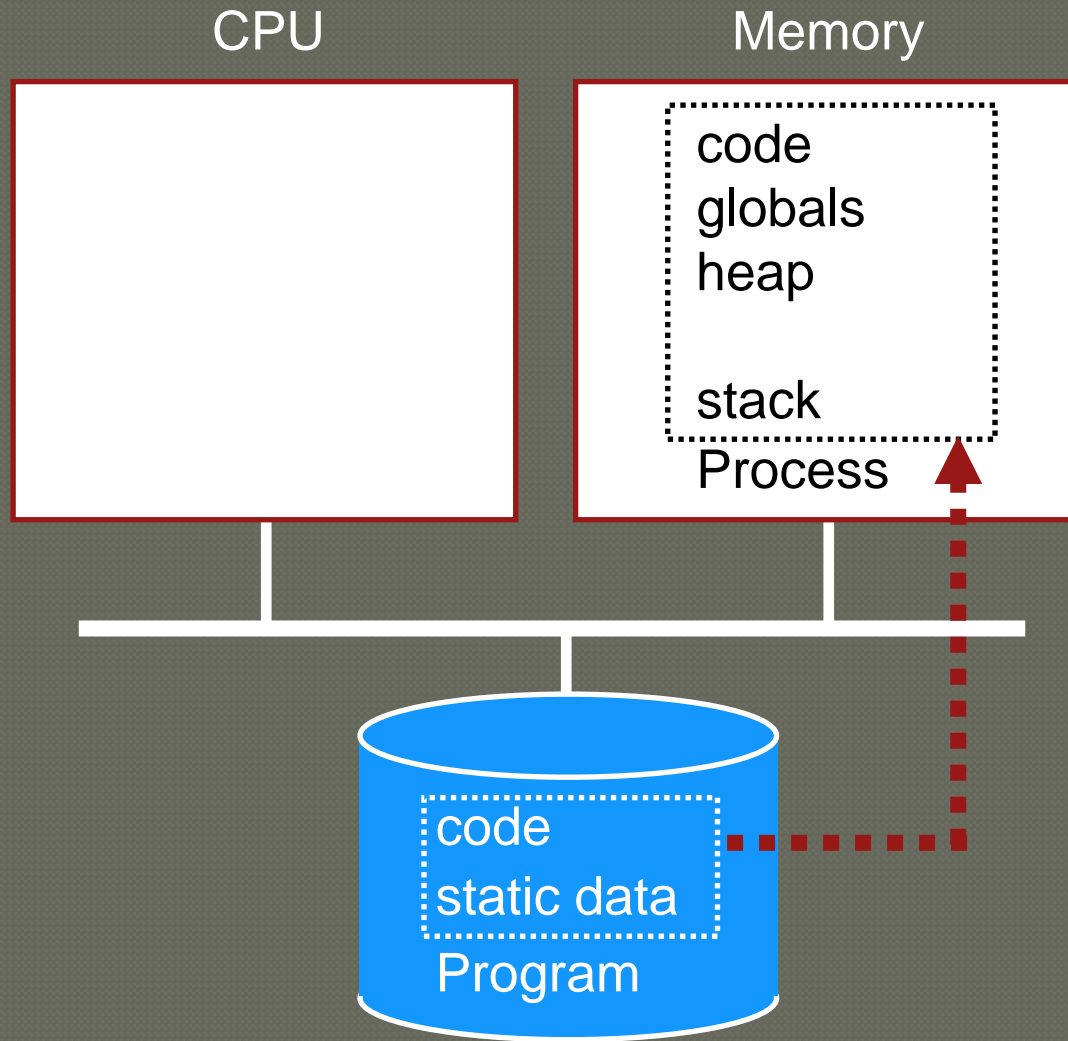
● Control

- Process creation
 - What does OS do when a process is created?
 - assigns a new unique ID
 - allocates space for the process in memory
 - initializes its process control block & sets it in place (e.g. in process list)

Process Creation



Process Creation



Control blocks
States
Description
Control

Processes

Dispatch Mechanism

Process is running- how to switch to other process?

Q1: How does Dispatcher get CONTROL?

Option 1: Cooperative Multi-tasking

- Trust process to relinquish CPU to OS through traps
 - Examples: System call, page fault (access page not in main memory), or error (illegal instruction or divide by zero)
 - Provide special `yield()` system call


Cooperative Approach

P1




yield() call

Cooperative Approach



Cooperative Approach



Cooperative Approach

P2

↑ yield() return

Cooperative Approach

P2



yield() call

Processes

Q1: How does Dispatcher get CONTROL?


- Problem with cooperative approach? **YES**
- Disadvantages: Processes can misbehave
 - By avoiding all traps and performing no I/O, can take over entire machine
 - Only solution: Reboot (windows 95)!
- **Not performed in modern operating systems**

Q1: How does Dispatcher get CONTROL?

Option 2: True Multi-tasking

- Guarantee OS can obtain control periodically
- Enter OS by enabling periodic alarm clock
 - Hardware generates timer interrupt (CPU or separate chip)
 - Example: Every 10ms
- User must not be able to mask timer interrupt
- Dispatcher counts interrupts between context switches
 - Example: Waiting 20 timer ticks gives 200 ms time slice
 - Common time slices range from 10 ms to 200 ms

Interrupts



Topics

● Everything about Processes

- Elements
- Control blocks
- States
- Description
- Control

● OS Execution



Done!